

Spring AI: A Field Guide

Spring AI brings the Spring ecosystem's familiar design principles to AI applications, offering one portable API in front of many providers.

Retrieval Augmented Generation (RAG) combines a vector store with a chat model so answers are grounded in your own documents instead of only the model's training data.

The ETL pipeline that feeds a vector store has three stages: Extract, via a `DocumentReader` such as `PagePdfDocumentReader`; Transform, via a `DocumentTransformer` such as `TokenTextSplitter`, which splits long text into smaller chunks sized for embedding; and Load, via a `DocumentWriter`, most commonly a `VectorStore`, which embeds each chunk and persists it.

Once documents are loaded, a `RetrievalAugmentationAdvisor` can retrieve the nearest chunks for a user question and fold them into the prompt before the chat model answers.